

Building Your First Amazon Redshift Data Warehouse

ANSWERING THE QUESTION, 'WHY AMAZON REDSHIFT'?



Russ Thomas

DATA ARCHITECT

@sqljudo www.sqljudo.com



Course Assumptions



My best guess about your objectives. Was I close?



Some background with relational and OLAP will be helpful



We will be moving along rapidly



Overview



Redshift History

Redshift Architecture

Redshift Potential






ABACOS WIDGETS



Abacos Data Warehouse Needs



Massive size potential



Scale out capability



Compatibility with existing reporting and BI tools



Abacos Data Warehouse Decision



Amazon Redshift is our go-forward technology



This module will focus on why



Remaining modules will focus on how



What's in It for You



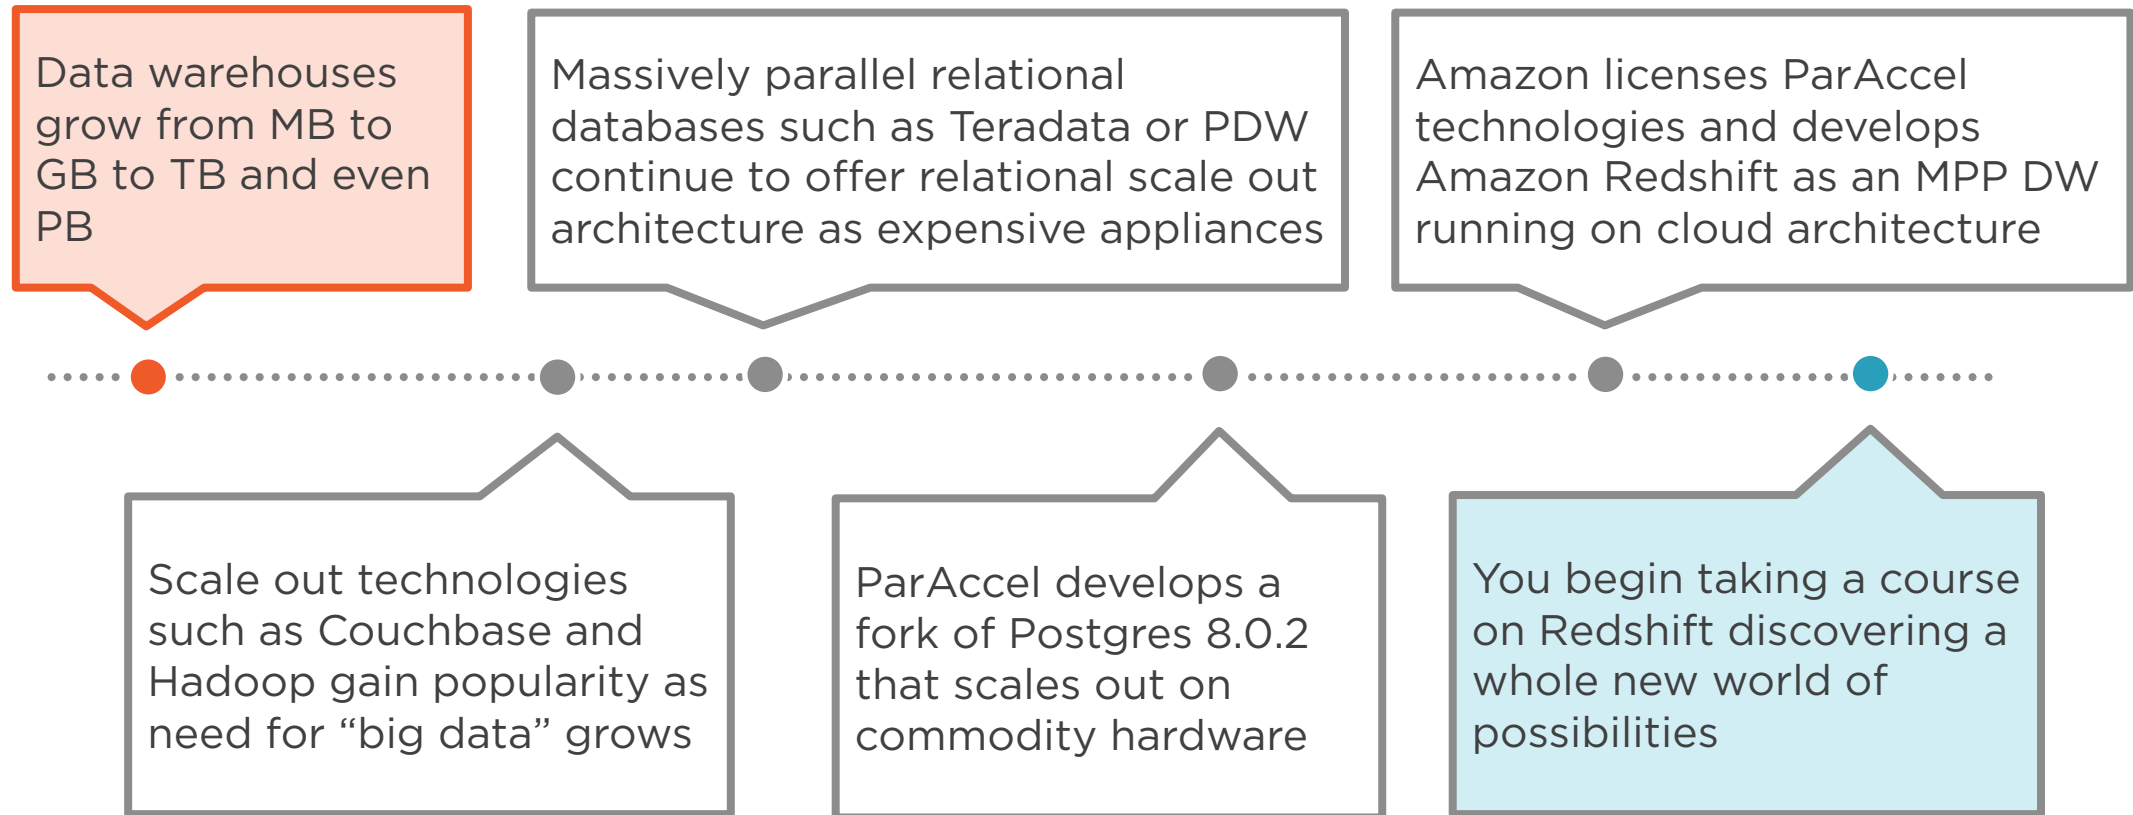
Experience building a fully functional DW on Redshift




Low to no cost trial



History of Amazon Redshift



Things to Remember About Amazon Redshift



Based on Postgres 8.0.2 but very much it's own technology now



Not everything supported in Postgres is supported in Redshift



Nerd note about a “Redshift”



Summary



Amazon Redshift:


- Massively parallel
- Relational (SQL) based
- Highly elastic cloud infrastructure








Redshift Cluster

Leader Node



- Client interaction
- Execution plan \ optimization
- Instruction compilation
- Work distribution



Compute Node(s)

- Dedicated OS, CPU, memory, storage
- Node scalability
- Dense storage \ dense compute
- Node slices



Redshift Query Optimization



**Query
optimization**



**Internal
statistics and
heuristics**



**ACID
compliance**



**Data
distribution**



Distribution Styles



Join performed



The goal of selecting a distribution style is minimizing impact of redistribution step




Efficient Distribution Join



Join performed



Data Warehouse Schema



Distribution Styles



Even distribution



Key distribution




All distribution





Even Distribution

- Distributes values evenly
- Default distribution style
- Good to average performance



Key Distribution


- Rows distributed by selected key
- Known relationships distributed across the same nodes
- Very effective for minimizing row redistribution costs




All Distribution

- All rows are copied to all nodes
- Significant increase to storage and cost of insert, updates, and deletes
- Best for medium size tables that change rarely and actively joined


Data Warehouse Schema



Redshift Performance Features




Row-based Architecture



Smaller storage
chunks often called
pages

Employee	Name	Social Security #	Address	Salary
----------	------	-------------------	---------	--------

Row-based Architecture



Smaller storage
chunks often called
pages



Pages contain all or
most columns from
each row



Reading large number
of column values is
inefficient

Row-oriented Storage

Employee	Name	Social Security #	Address	Salary
----------	------	-------------------	---------	--------




Column-oriented Storage

Employee	Name	Social Security #	Address	Salary
Employee	Name	Social Security #	Address	Salary
Employee	Name	Social Security #	Address	Salary




Row-based Architecture


Employee	Name	Social Security #	Address	Salary
----------	------	-------------------	---------	--------




Row-based Architecture




Row-based Architecture



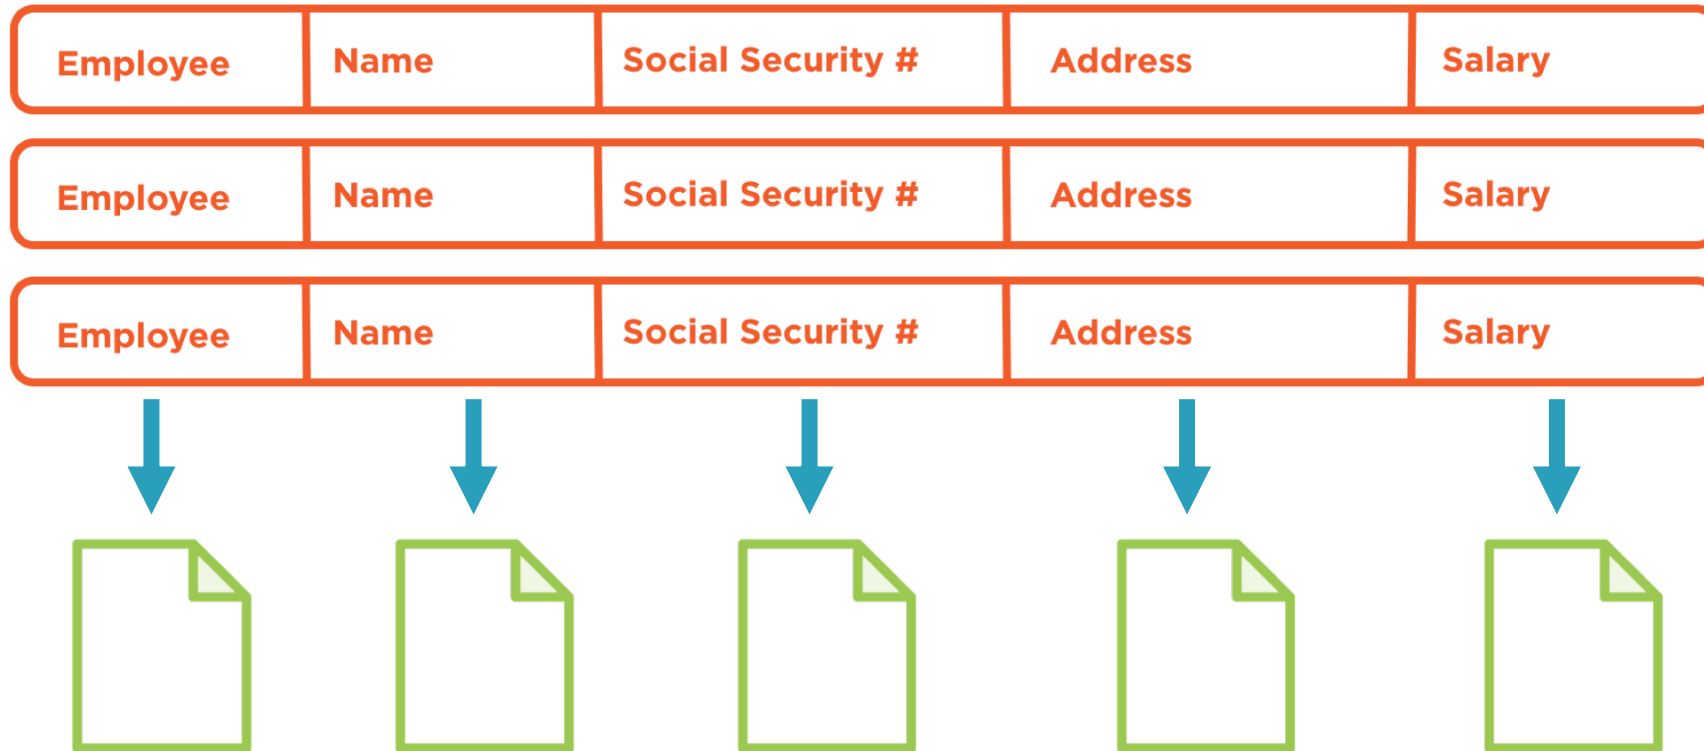
Row-based Architecture



Row-based Architecture



Column-based Architecture



Column-based Architecture



There is no clustered or non-clustered index



There is column-oriented, massively parallel, scale out architecture



Not effective for large amount of CRUD or small searches



AWS Database Offerings

- Relational database services
- NoSQL services
- Petabyte scale data warehousing






Amazon Redshift Sorting

- A sortkey sorts data on disk
- Declared like distkey and primary key
- Part of create statement
- Only one sortkey permitted

Column-based Architecture





Effective Data Compression

- Optimized for single data type
- Speeds data reads
- Reduces overhead of redistribution
- Customizable compression

Amazon Redshift Data Warehouse



Designed with 100% OLAP focus



Most limitations are from an OLTP perspective only



Homework

- Review full schema creation scripts
- Identify different approaches
- Ensure sortkeys, distkeys make sense

Next up: loading our data warehouse

