

Populating Redshift



Russ Thomas

DATA ARCHITECT

@sqljudo www.sqljudo.com



Module Prerequisites



Redshift Cluster with empty table schema




Variances in Distribution Keys / Sort Keys is ok




Access to download content for the course




The Copy Command




Logical
approach



Data formats




Command
syntax




Source locations



Integrating Redshift into Your ETL




Leveraging
Copy



Staging Tables




Transformations



Automated
Loads



Leveraging Data Streaming



Loading into Nodes and Node Slices




ETL Approaches to Consider



Upserts



Overhead of major update / insert activity



Traditional ETL approach

Let each of your tools do
what they do best!





A scripting approach does not preclude:

- Parallel tasks
- Multiple files
- Advanced approaches through SDK

The Copy command is the de facto standard for bulk loading data and ongoing population of your cluster



Overhead of Major Update / Insert Activity



Updates act as deletes followed by an insert



Inserts are not automatically sorted per sort key



Deletes are marked as deleted but continue to use space on disk



Vacuum Process

- Resorts rows per the sort key
- Reclaims space from deleted rows
- Includes option to re-index
- Substantially different from Postgres vacuum process

Copy Command



- Sources include:
 - S3 storage
 - External SSH sources
 - AWS databases such as DynamoDB
 - Kinesis Firehose
- Authorization
 - IAM Roles
 - Access Keys
- Options include:
 - GZIP and other zip technologies
 - Multiple formats and delimiters
 - Manifest files for parallel loads

